

Regional LLMs by Continent

Executive summary

This report ranks the **top ten LLMs per continent by suitability for understanding and generating the major native languages used in that region**, not by raw general intelligence alone. I weighted models toward **explicit regional language support and regional corpus adaptation** over generic frontier capability, because that better matches the user’s objective. The scoring rubric used throughout was:

Regional-language fit 40%; regional corpus or continued pretraining 20%; benchmark evidence on multilingual or region-specific evaluations 15%; general generation quality 10%; deployability and access 10%; price/transparency 5%. Where a continent has a thin sovereign-LLM ecosystem, I ranked the **best available multilingual substitutes** but penalised them for weak or undocumented support for the region’s low-resource languages. This matters especially for Africa, South America, Oceania, and North America’s Indigenous languages. Recent benchmark literature also shows that low-resource and culturally grounded evaluation remains a major gap across continents. ¹

The headline finding is unevenness. **Asia** now has the deepest regional ecosystem, with serious sovereign or region-first model work spanning Chinese, Indic, Southeast Asian, Arabic, Korean, and Japanese language clusters. **Europe** is the next strongest, especially for “digital sovereignty” efforts such as EuroLLM, Lucie, Pharia, Poro, CroissantLLM, and BgGPT. **South America** has made notable progress with **Latam-GPT**, **Sabiá**, and **Tucano**, but Indigenous-language support is still early. **Africa** has genuine breakthroughs—especially **Mansa**, **InkubaLM**, and **Jacaranda’s African-language Llama derivatives**—but many are still narrow, translation-heavy, or limited in breadth. **North America** paradoxically has excellent English/French/Spanish coverage via general frontier models yet very weak officially documented support for most Indigenous North American languages. **Oceania** is even thinner: from the primary sources reviewed, there is no mature, released, general-purpose sovereign Māori or Pacific-islands LLM comparable to the best Asian or European regional stacks. **Antarctica** is a special case: because it has no native human population, the right ranking criterion is multilingual station interoperability and offline deployability rather than “regional nativeness.” ²


A compact readiness heatmap is below. Darker means stronger evidence from primary sources and post-2022 benchmarks.

Continent	Regional-model depth	Benchmark maturity	Open-weight options	Indigenous / low-resource readiness
Africa	High	Medium	High	Medium
Asia	Very high	High	Very high	Medium
Europe	Very high	Medium	Very high	Low-medium
North America	Low	Low	High	Low

Continent	Regional-model depth	Benchmark maturity	Open-weight options	Indigenous / low-resource readiness
South America	Medium	Medium	Medium-high	Low-medium
Oceania	Low	Low	Medium	Low
Antarctica	N/A	N/A	High	N/A

This matrix reflects the concentration of region-specific releases and benchmarks such as **AfroBench**, **IrokoBench**, **MILU**, **IndicMMLU-Pro**, **SEA-LION/SEA-HELM work**, and **Latam-GPT's Trueque and Choclo**, alongside official model documentation for open and commercial models. ³

The release cadence shows how quickly the regional-LLM landscape has matured since 2023. ⁴



Open questions / limitations: some providers do not publicly disclose model sizes, full language lists, or current API prices; some contemporary model pages have already shifted to successor models; and for North American Indigenous and Pacific languages in particular, the public evidence base is too thin to rank many “true regional” LLMs with high confidence. Where this happened, I say so explicitly rather than guessing. ⁵

Africa

Africa now has a recognisable regional-LLM stack, but it is still fragmented. The strongest Africa-first options are either **narrow but high-fit**—for example Swahili-, Yoruba-, or isiXhosa/isiZulu-adapted models—or **enterprise African-language platforms** such as Mansa. Benchmark work such as

AfroBench and **IrokoBench** also shows that many large generalist models degrade sharply on African-language reasoning and culturally grounded tasks, which is why this ranking heavily favours explicit African-language adaptation over generic “multilingual” claims. ⁶

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations
1	Mansa — African Languages Lab, large-scale multimodal foundation model, 2025–2026	Broadest Africa-first production fit in the sources reviewed	30–40 African languages; 19B tokens of proprietary African text; text, speech, and vision; languages include Hausa, Yoruba, Swahili, Igbo, Amharic, Twi, Kinyarwanda, Zulu and more	Integrated into Phrase BYOE and Blackbird; marketed as enterprise African-language AI	Cloud/API/portal; proprietary; pricing not publicly enumerated on the pages reviewed	More platform-like than a downloadable open LLM; public benchmark detail is still thinner than for research models
2	InkubaLM 0.4B — Lelapa AI, 400M, 2024	Best-documented small Africa-first LLM for edge or low-compute contexts	Designed for low-resource African languages; balanced across Swahili, Hausa, Yoruba, isiXhosa, isiZulu	Paper reports strong task performance and consistency in those five languages	Small open model; local deployment attractive for on-device or low-connectivity use	Small parameter count constrains broad reasoning and long-form generation
3	AfroLlama_V1 — Jacaranda, 8B, 2024	Best open multi-language African chat model in the reviewed sources	Swahili, Hausa, Yoruba, isiXhosa, isiZulu, English; built by pretraining and merging African-language Llama 3 variants	Publicly released via model cards; created to support Jacaranda’s multilingual health and support workflows	Open weights on Hugging Face; local/self-hosted	Language set is only five African languages; still narrower than continent scale

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations
4	UlizaLlama3 — Jacaranda, 8B, 2024	Strongest Swahili-specific Africa LLM among the primary sources gathered	Swahili-focused continual pretraining and tokenizer adaptation	Originated in maternal-health support; Jacaranda positions it for AI-driven support at scale	Open weights; self-hostable	Swahili-first rather than pan-African
5	YorubaLlama — Jacaranda, 8B, 2024	Best clearly documented Yoruba-generative candidate in the reviewed primary sources	Enhanced to generate and understand Yoruba	Released as part of Jacaranda's five-language African expansion	Open weights; self-hostable	Narrow language scope; less public benchmark evidence than broader African evaluations
6	Xhosa_ZuluLlama3_v1 — Jacaranda, 8B, 2024	Valuable for southern African deployment fit	Optimised for isiXhosa and isiZulu understanding and generation	Released as part of Jacaranda's five-language suite	Open weights; self-hostable	Narrow linguistic focus; public evaluation detail limited
7	Jais 30B Chat — Inception / G42 / Cerebras, 30B, 2023	Best-documented large Arabic-first LLM relevant to North Africa	Arabic and English; trained on Arabic- and English-language corpora	One of the earliest major Arabic-centric instruct LLMs	Open weights via model card; self-hostable; commercial licence terms attached	Strong for Arabic, but not for Sub-Saharan African languages
8	SILMA 9B Instruct — MBZUAI, 9B, 2024	Important for Maghrebi / dialectal Arabic, which generic Arabic models often underserve	Arabic dialect focus, including low-resource dialectal varieties	Built to improve dialectal Arabic modelling beyond MSA-centric systems	Open weights via Hugging Face	Narrow regional coverage outside Arabophone Africa; dialectal focus not continent-wide

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations
9	Aya Expanse 32B — Cohere For AI, 32B, 2024	Best open multilingual fallback when African-specific models are unavailable	23 optimised languages including Arabic; broader Aya work targets multilingual inclusion at scale	Outperforms many open rivals on multilingual evaluations; useful where Arabic is required alongside global languages	Open weights; CC-BY-NC research licence; self-hostable	Official optimised list does not explicitly include most African low-resource languages such as Yoruba or isiZulu
10	Gemini 2.5 Pro — Google, proprietary, 2025	Best commercial fallback for broad multilingual and multimodal African deployments	Google documents text output in 140+ languages; Live API list explicitly includes African languages such as Hausa, Swahili, Somali, Yoruba and Zulu	Strong cloud production option; wide modality support	Cloud only; public pricing available on Vertex/ Gemini API	Not Africa-trained; “support” does not imply Africa-specific cultural grounding

Africa verdict: if you need the **best Africa-first fit**, use **Mansa, InkubaLM**, or **Jacaranda’s African-language family**. If you need **North African Arabic**, use **Jais** or **SILMA**. If you need a strong **commercial multilingual fallback**, use **Gemini 2.5 Pro** but expect noticeably weaker cultural grounding for many low-resource African languages than the region-first models. AfroBench and IrokoBench strongly support that caution. ¹⁷

Asia

Asia has the strongest and most diverse regional LLM ecosystem in this review. The key reason is coverage across multiple language macro-regions: **Chinese, Indic, Southeast Asian, Arabic, Korean**, and **Japanese**. Unlike other continents, Asia now has several models explicitly trained on regional corpora and tested on regional benchmarks such as **MILU, IndicMMLU-Pro**, and **SEA-LION / SEA-HELM / SeaBench-related evaluations**. ¹⁸

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Source
1	Qwen3 32B — Alibaba, 32B dense, 2025	Broadest explicit Asian-language coverage among open models in the reviewed sources	119 languages/dialects; includes Chinese varieties, Arabic dialects, Indic languages, Thai, Lao, Japanese, Korean, Indonesian, Malay, Tagalog, Cebuano, and Tok Pisin	Official blog claims competitive frontier-level performance; huge pretraining set of ~36T tokens	Apache 2.0; open weights; local deployment via vLLM, SGLang, Ollama, LM Studio, llama.cpp	Not specifically “Asia-sovereign” in every subregion; hosted pricing varies by provider	19
2	SEA-LION v3 — AI Singapore-led consortium, 8B/9B instruct, 2025	Best documented SEA-native family	Supports 11 SEA languages: English, Chinese, Indonesian, Vietnamese, Malay, Thai, Burmese, Lao, Filipino, Tamil, Khmer	Paper reports state-of-the-art performance for SEA-supported LLMs	Open source/open models	Stronger in Southeast Asia than in East, West, or South Asia	20
3	Sailor2 20B — Sea AI Lab, 20B, 2025	One of the strongest SEA-centric open families	Family designed for 15 Southeast Asian languages	Frequently cited in SEA evaluations as a leading regional open model family	Open weights; local deployment	Focused on SEA, so weaker fit for Arabic, Korean, or Japanese	21
4	Jais 30B Chat — Inception / G42 / Cerebras, 30B, 2023	Strongest large Arabic-first generative model in the reviewed Asia sources	Arabic + English; trained on Arabic/English corpora	Important baseline for West Asia / Gulf deployments	Open weights with licence restrictions	Limited for non-Arab Asian language families	13
5	HyperCLOVA X — NAVER, proprietary, 2024	Best Korean-first commercial model in the reviewed set	Korean language and culture; multilingual ability reported for Korean↔English/Japanese/Chinese	Technical report says Korean reasoning and cultural knowledge are standout strengths	Cloud / proprietary	Closed model and relatively limited public pricing detail	22

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Source
6	CyberAgent CALM3 22B Chat — CyberAgent, 22B, 2024	Strong Japanese open model with substantial pretraining	Pretrained from scratch on 2T tokens; Japanese-specialised chat model	Strong fit for Japanese enterprise and consumer NLP	Open model via Hugging Face	Primarily Japanese; weaker outside Japan compared with Qwen3 or Gemini	23
7	Krutrim-2 Instruct — OLA Krutrim, 12B, 2025	Strong India-first commercial/open-weight candidate	English + Indian languages; trained on code, math, Indic languages, Indian-context data, books, synthetic data	Designed for Indian demographic and multilingual workloads	Model card available; self-hosting possible; enterprise offering	Public benchmark transparency is thinner than for MILU-style academic evaluations	24
8	OpenHathi-7B-Hi — Sarvam AI, 7B, 2023	Good low-cost Hindi / Hinglish fit	Hindi, English, Hinglish; first OpenHathi release	Positioned as a frugal Hindi-first base model	Open weights; easy local serving via vLLM / Ollama	Hindi-centric; ageing relative to newer 2025–2026 families	25
9	Llama-3 Nanda 10B Chat — MBZUAI, 10B, 2025	Strong bilingual Hindi-English specialist	Trained on 65B Hindi tokens; bilingual Hindi/English	Useful for high-Hindi workloads where a compact open model is desired	Open weights via model card	Narrower than Krutrim or Qwen3 for multilingual India	26
10	Gemini 2.5 Pro — Google, proprietary, 2025	Best commercial all-Asia breadth fallback	Google documents 140+ languages; explicit Live API support spans Arabic, Hindi, Bengali, Gujarati, Kannada, Malayalam, Marathi, Nepali, Odia, Punjabi, Tamil, Telugu, Thai, Vietnamese, Japanese, Korean and more	Deep multimodal and production strength	Cloud only; public pricing available	Not region-trained; may underperform region-tuned models on culture-heavy local tasks	16

Asia verdict: for **maximum breadth**, choose **Qwen3 32B**; for **Southeast Asia**, choose **SEA-LION v3** or **Sailor2**; for **Arabic West Asia**, choose **Jais**; for **Korean**, **HyperCLOVA X**; for **Japanese**, **CALM3**; for **Indic**, choose **Krutrim-2**, **OpenHathi**, or **Nanda** depending on whether you want breadth, openness, or Hindi intensity. ²⁷

Europe

Europe’s strongest pattern is **sovereign multilinguality**. Several projects are explicitly motivated by European language coverage, legal compliance, or cultural rights. Because Europe has many high-resource languages and public institutions, its best models tend to be **open-weight multilingual foundations** rather than monolingual low-resource specialists. That said, the best sovereign European models still focus disproportionately on the large Western European languages, with comparatively less public evidence for smaller and minority languages. ²⁸

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
1	EuroLLM — EuroLLM consortium, open models, 2024	Best pan-European fit in the reviewed sources	Optimised for all 24 official EU languages plus additional European languages	Explicitly designed as a multilingual European foundation model	Open weights; self-hostable	Public commercial deployment evidence still developing	²⁹
2	Lucie-7B Instruct v1.1 — OpenLLM-France, 7B, 2025	Strongest genuinely open, rights-conscious European multilingual model in this review	3T-token multilingual dataset centred on French, plus English, German, Spanish, Italian and parallel data	Paper positions Lucie as one of the first OSI-compliant language models	Open source/ open weights; GGUF variants available	Much stronger on core Western European languages than on the full EU long tail	³⁰
3	CroissantLLM Chat — CroissantLLM, 2024	Strong Europe-first multilingual open model	European-language focus, especially French-centred multilinguality	Recognised as a European multilingual model family	Open weights	Public evidence base narrower than EuroLLM or Lucie	³¹

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
4	Pharia-1-LLM-7B-control — Aleph Alpha, 7B, 2024	Strong sovereign-EU compliance option	Multilingual corpus; linguistically and culturally optimised for German, French, Spanish; trained with EU compliance emphasis	Positioned for regulation-sensitive European deployments	Open Aleph licence; on-prem possible in enterprise contexts	Non-commercial research licence is more restrictive than Apache 2.0-style models	32
5	Poro 34B Chat — Silo AI / TurkuNLP / HPLT, 34B, 2024	Best strong Finnish open model in this review	Finnish, English, code; 1T tokens	Paper and model card position it as best-in-class for Finnish at release	Apache 2.0; open weights	Narrower linguistic scope than EuroLLM or Lucie	33
6	BgGPT 3.0 — INSAIT, Gemma-3-based 4B/12B/27B, 2026	Best Bulgaria-specific modern release found	Bulgarian-adapted series based on Gemma 3	Evaluated on Bulgarian-translated and Bulgarian-specific benchmarks	Open on Hugging Face; GGUF variants available	Mainly Bulgarian-focused, not pan-European	34
7	Claire-7B-FR-Instruct — OpenLLM-France, 7B, 2026	Strong French conversational specialist	French-language instruct-tuned model descended from Claire family	Useful for French dialogue and instruction following	Open weights	Primarily French; less suitable for cross-Europe multilingual workloads	35
8	Gemma 3 27B — Google, 27B, 2025	Best open multilingual fallback with local deployment	140+ languages; runs locally on consumer-grade GPUs with QAT	Strong open multilingual utility when sovereign models are unavailable	Open weights under Google usage terms; self-host or cloud	Not Europe-trained and less culturally grounded than EuroLLM / Lucie / Pharia	36

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
9	Command A — Cohere, proprietary/open-servable model, 2025	Strong commercial EU-language enterprise choice	23 languages explicitly include English, French, Spanish, Italian, German, Polish, Czech, Romanian, Greek, Dutch, Ukrainian and more	Enterprise agents, RAG, citations, tool use	Cloud and deployable on two A100/H100s; current public API price \$2.5 in / \$10 out per 1M	Not Europe-specific and language list omits some smaller European languages	37
10	Llama 3.3 70B — Meta, 70B, 2024	Strong open fallback for major Western and Southern European languages	Official support across English plus French, German, Italian, Portuguese, Spanish, and Hindi/Thai	Widely deployable and fine-tunable	Open weights under Llama community licence	Officially supported language set is narrower than “Europe-wide” claims often made by users	38

Europe verdict: if you want **highest European-language suitability**, choose **EuroLLM** first, then **Lucie**, then **CroissantLLM** or **Pharia** depending on whether you value openness or compliance. For country-specific specialisation, **Poro**, **BgGPT**, and **Claire** are the standouts. [39](#)

North America

North America is the hardest inhabited continent to rank if the criterion is **native-language justice** rather than raw quality. For English, French, and Spanish, North American users are exceptionally well served by general frontier models. For **Indigenous North American languages**—for example Navajo, Cherokee, Cree, Ojibwe, Inuktitut, and many others—the public evidence base is much weaker. In the primary sources reviewed, I did **not** find a mature, released, general-purpose LLM with strong official support for those languages comparable to the best Africa-, Asia-, or Europe-specific systems.

Brookings’ 2025 review of small language models for Indigenous-language revitalisation helps explain why: low-resource, community-governed model work is still at an early stage. [40](#)

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
1	Command A — Cohere, 2025	Best explicit language list for the main North American public languages	English, French, Spanish explicitly supported at high quality; enterprise-oriented	Strong for RAG, cross-lingual tasks, long context	Cloud or self-deploy; \$2.5 in / \$10 out per 1M	No published official support for most Indigenous North American languages	37
2	GPT-4.1 — OpenAI, 2025	Very strong English-first, high-quality North American fallback	OpenAI says models are multilingual but still optimised for English	Excellent general generation and tool use	Cloud only; API page shows roughly \$2 in / \$8 out per 1M	Official language documentation is broad rather than explicit; Indigenous support undocumented	41
3	Claude Sonnet 4 — Anthropic, 2025	Top-tier English enterprise option with strong French/Spanish use in practice	Broad multilingual use, though official docs are less explicit than Cohere's	Strong reasoning and enterprise fit	Cloud/API; \$3 in / \$15 out per 1M	Weak public documentation for low-resource regional languages	42
4	Gemini 2.5 Pro — Google, 2025	Best breadth and multimodality	140+ languages; broad cloud production support	Strong option across English, French, Spanish and mixed-media workflows	Cloud only; Vertex public pricing available	Still not Indigenous-language-first	43
5	Aya Expansive 32B — Cohere For AI, 2024	Best open multilingual alternative with explicit English/French/Spanish support	23 optimised languages include English, French, Spanish, Portuguese	Strong multilingual benchmark performance	Open weights; CC-BY-NC research licence	No explicit support documentation for most Indigenous North American languages	15

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
6	Gemma 3 27B — Google, 2025	Best open local model for North American deployments needing privacy	140+ languages, local/open deployment	Good for laptops, desktops, and private cloud	Open weights; local friendly	Region fit is generic, not North America-specific	36
7	Llama 3.3 70B — Meta, 2024	Strong open fallback for English, French, Spanish, Portuguese	Explicitly supports English plus major European languages including French, Spanish, Portuguese	Fine-tune and run anywhere	Open weights under Llama community licence	Official supported-language set is relatively narrow; no Indigenous-language guarantees	38
8	Qwen3 32B — Alibaba, 2025	Very broad multilingual open model, useful for immigrant-language-rich settings	119 languages/dialects including Spanish and Haitian	Strong multilingual breadth	Apache 2.0; local deploy	Not trained for North American Indigenous languages specifically	44
9	Mistral Small 3.1 — Mistral AI, 24B, 2025	Good multilingual open commercial option for English/French/Spanish workloads	Dozens of languages including French, Spanish, Portuguese	Strong small-model multilinguality	Open weights; pricing for this exact retired model is less prominent on current pages	Not region-specific; current public pricing pages emphasise successors	45
10	Gemini Live / Native Audio family — Google, 2025–2026	Included because speech-first multilingual support often matters in North America	Official list covers 97 Live API languages; useful for multilingual voice interfaces	Strong for realtime voice and translation workflows	Cloud only	Still not Indigenous-North-America-first, though it does broaden accessibility	46

North America verdict: for mainstream public-language coverage, the best stack is **Command A**, **GPT-4.1**, **Claude Sonnet 4**, and **Gemini 2.5 Pro**. But if your actual goal is **Indigenous-language**

revitalisation, the public frontier ecosystem is still inadequate; community-governed small language models are likely the more realistic path than generic frontier LLMs. ⁴⁷

South America

South America is now the clearest example of a continent moving from “generalist dependency” toward **regional AI sovereignty**. The major story is the emergence of **Latam-GPT** plus a strong Brazilian Portuguese sub-ecosystem around **Sabiá** and **Tucano**. The shortcoming is that **Spanish and Portuguese are much better served than Quechua, Aymara, Guaraní, or other Indigenous South American languages**, although Latam-GPT explicitly frames Indigenous-language inclusion as a future direction. ⁴⁸

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
1	Latam-GPT — CENIA-led coalition, 2026	Best continent-level regional fit	Open regional model built from Latin American data and identity; Spanish and Portuguese first; plans for Indigenous languages including Rapa Nui and others	Backed by newly introduced region-specific benchmarks Trueque and Choclo ; collaboration across multiple Latin American institutions	Open/ regional project; pricing not applicable like an API-first proprietary model	Still early-generation; public consumer-grade benchmark breadth remains limited versus frontier cloud models	⁴⁹
2	Sabiá-4 — Maritaca AI, 2026	Best Brazilian Portuguese commercial model in the reviewed sources	Portuguese and Brazilian context; four-stage pipeline with Portuguese and Brazilian legal corpora, long context, instruction tuning, alignment	Sabiá family is positioned as high-fidelity Portuguese chat for Brazil	API/cloud only by default; public price R\$5 input / R\$20 output per 1M	Portuguese-first rather than South America-wide; weights not generally downloadable	⁵⁰

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
3	Tucano 2 — Polyglot / N. Klugé Corrêa, 0.5B–3.5B, 2026	Best fully open Portuguese sovereign family	Native Portuguese pretraining; open Portuguese models that outperform prior Portuguese models of similar size	Paper and collection explicitly position Tucano 2 as a better open Portuguese stack	Apache 2.0; self-hostable	Smaller than frontier models; Portuguese-centric	51
4	Gemini 2.5 Pro — Google, 2025	Best commercial breadth fallback for Spanish/Portuguese workflows	140+ languages; production-grade multilinguality	Strong multimodal cloud option	Cloud only; public pricing available	Less Latin American cultural grounding than Latam-GPT or Sabiá/Tucano	43
5	Command A — Cohere, 2025	Best commercial structured-output and RAG fit for Spanish + Portuguese	Explicit support for Spanish and Portuguese plus major European languages	Enterprise RAG and agent use cases	Cloud or self-deploy; \$2.5 in / \$10 out per 1M	No Indigenous South American focus	37
6	Aya Expansive 32B — Cohere For AI, 2024	Strong open multilingual option for Spanish and Portuguese	Explicitly optimised for Portuguese and Spanish among 23 languages	Strong multilingual benchmark performance	Open weights; research licence	Not Latin-America-specific	15
7	Qwen3 32B — Alibaba, 2025	Broad multilingual open fallback	Supports Portuguese and Spanish plus Haitian, Papiamentu, Kabuverdianu and other “Other” languages in the official list	Good multilingual reasoning breadth	Apache 2.0; local	No Latin American corpus specialisation	44

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
8	Gemma 3 27B — Google, 2025	Best local open model to pair with regional data	140+ languages and local deployability	Strong practical open deployment for South American institutions	Open weights; local-friendly	Generic multilingual rather than region-first	52
9	GPT-4.1 — OpenAI, 2025	Excellent general text quality for Spanish and Portuguese	Multilingual, though English-optimised	Strong fallback when quality is prioritised over openness	Cloud only; about \$2 / \$8 per 1M	Weak public documentation for Indigenous-language support	41
10	Claude Sonnet 4 — Anthropic, 2025	Strong closed-model fallback for Spanish/Portuguese enterprise use	Broad multilingual use	Enterprise-grade reasoning and writing	Cloud only; \$3 / \$15 per 1M	Not region-specific; low transparency on exact language performance	42

South America verdict: if the priority is **regional language authenticity**, the order is clear: **Latam-GPT**, then **Sabiá-4**, then **Tucano 2**. If the priority is more general production capability across Spanish and Portuguese, use **Gemini**, **Command A**, or **Aya Expanse**. 53

Oceania

Oceania has the weakest released general-purpose regional-LLM ecosystem among the inhabited continents reviewed here. The two strongest explicit pieces of evidence I found were: **Google's official Māori support in Gemini Live**, and **Qwen3's official Tok Pisin support**. Beyond that, the public ecosystem is notably thin. A 2025 plan for a sovereign Māori language model suggests the field is still in formation, while Te Hiku Media's strongest publicly visible AI milestone remains a speech model rather than a general-purpose LLM. For this reason, Oceania rankings here are pragmatic: they favour models with the *best currently documented fit* for English plus a small number of explicitly documented Oceanian languages. 54

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
1	Gemini 2.5 Pro / Gemini Live family — Google, 2025–2026	Best explicit Oceania fit because Māori is officially listed	140+ languages; Live API list explicitly includes Maori	Strong multimodal and voice workflows	Cloud only; public pricing available	No explicit Samoan, Tongan, or Tok Pisin in the Live API list reviewed	55
2	Qwen3 32B — Alibaba, 2025	Best open Oceania fit because Tok Pisin is explicitly listed	119 languages/ dialects; official list includes Tok Pisin	Strong open multilingual reasoning option	Apache 2.0; local deploy	No explicit Māori, Samoan, or Tongan in the official list reviewed	19
3	Gemma 3 27B — Google, 2025	Best local open option for Oceania institutions	140+ languages; strong local deployability	Useful for privacy-sensitive local deployment in Australia/ NZ/Pacific orgs	Open weights; local friendly	Publicly available language detail is broad, not Pacific-specific	36
4	Command A — Cohere, 2025	Strong English-first enterprise model with good Asia-Pacific business-language mix	English plus Indonesian and other APAC-relevant languages in the 23-language list	Strong for agents and cross-lingual enterprise use	Self-deploy or cloud; \$2.5 / \$10 per 1M	No documented Māori or Tok Pisin support in the official list	37
5	GPT-4.1 — OpenAI, 2025	Very strong English fallback for Australia/NZ workloads	Multilingual but English-optimised	Excellent generic generation	Cloud only; about \$2 / \$8 per 1M	No Pacific-language specificity in the reviewed primary docs	41
6	Claude Sonnet 4 — Anthropic, 2025	Strong English enterprise fallback	Broad multilingual use in cloud settings	High-quality writing and reasoning	Cloud only; \$3 / \$15 per 1M	Weak transparency for Pacific-language support	42

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
7	Llama 3.2 3B / 3.1 family — Meta, 2024	Good open edge model for English deployments	Official supported set includes English and several European languages; runs on edge/mobile	Strong when local deployment matters more than Pacific language depth	Open weights; select edge/mobile fit	No official Māori or Tok Pisin support in the model cards reviewed	56
8	Mistral Small 3.1 — Mistral AI, 2025	Good multilingual open fallback for English and wider APAC interactions	Supports dozens of languages including Indonesian, Malay, Vietnamese and others	Strong small open model class	Open weights	No documented Pacific-islands specialisation	45
9	Aya Expansive 32B — Cohere For AI, 2024	Useful open multilingual fallback for English plus some adjacent Austronesian languages	Official 23-language list includes Indonesian; strong cross-lingual behaviour	Strong multilingual benchmark evidence	Open weights; research licence	No explicit Māori, Tok Pisin, Samoan, or Tongan in optimised language list	15
10	SEA-LION v3 — AI Singapore-led consortium, 2025	Included as the best nearby Austronesian/SEA open family	Covers Indonesian, Malay, Filipino and other SEA languages that can be relevant in Melanesian/ West Pacific settings	Strong SEA-focused results	Open models	Not Oceania-specific and no explicit Māori/Tok Pisin	20

Oceania verdict: the only models with **clear explicit regional-language evidence** from the sources reviewed are **Gemini** for **Māori** and **Qwen3** for **Tok Pisin**. Everything else is effectively an English-first or broad multilingual compromise. 57

Antarctica

Antarctica is a special case. It has **no native human population**, and today's operational reality is a rotating multinational research and logistics community spread across **34 national Antarctic**

programmes. That means the right question is not “which Antarctic-native LLM exists?” but rather **which models best support multilingual scientific operations, translation, documentation, and local/offline deployment in a harsh, bandwidth-constrained environment.** The ranking below therefore rewards **multilingual breadth, open deployment, and field practicality** more than cultural nativeness. ⁵⁸

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
1	Command A — Cohere, 2025	Best balance of multilingual breadth + self-deployable enterprise serving	23 languages cover English, French, Spanish, German, Portuguese, Russian, Arabic, Chinese, Japanese, Korean and more	Strong agentic and RAG use for operations, manuals, logs	Cloud or self-deploy; \$2.5 / \$10 per 1M	Not optimised for fully offline edge/mobile deployments	³⁷
2	Qwen3 32B — Alibaba, 2025	Broadest open multilingual coverage, strong for mixed international stations	119 languages/ dialects; includes Russian, Chinese, Japanese, Korean, Spanish, Portuguese and more	Strong open model for multilingual ops	Apache 2.0; local deploy	Requires substantial local compute in dense 32B form	⁴⁴
3	Gemma 3 27B — Google, 2025	Best open local-deployment practical choice	140+ languages; rich open deployment ecosystem; QAT makes local serving easier	Suitable for private scientific or field workflows	Open weights; local on consumer-class hardware possible with QAT	Not trained for the specific document culture of Antarctic operations	⁵²
4	Aya Expansive 32B — Cohere For AI, 2024	Strong open multilingual alternative with good cross-lingual quality	23 optimised languages including major station languages	Strong benchmark evidence against open peers	Open weights; research licence	Optimised language set narrower than Qwen3 or Gemma 3	¹⁵

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
5	Gemini 2.5 Pro — Google, 2025	Best cloud multimodal option for remote science and multimodal docs	140+ languages; text, code, images, audio, video	Useful for scientific multimodal workflows	Cloud only; public pricing available	Connectivity dependence is a major drawback for isolated sites	59
6	Mistral Small 3.1 — Mistral AI, 2025	Strong open-and-small compromise	Supports dozens of languages and multimodal inputs	Good for local or semi-local deployment if successor pricing/ model governance is acceptable	Open weights	Current public pages focus on successors; exact live pricing less straightforward	45
7	Llama 3.1 / 3.3 family — Meta, 2024	Strong open fallback with many mature serving options	Official support for English, French, German, Italian, Portuguese, Spanish, Thai, Hindi	Fine-tuneable and standardised infra support	Open weights under community licence	Official language set omits Russian, Chinese, Japanese, Korean from “supported” list	60
8	GPT-4.1 — OpenAI, 2025	High-quality cloud fallback for English- and Spanish-heavy teams	Multilingual but English-optimised	Strong for reports, coding, and document workflows	Cloud only; about \$2 / \$8 per 1M	Cloud dependence; no local/offline option	41
9	Claude Sonnet 4 — Anthropic, 2025	Strong cloud writing/ reasoning fallback	Broad multilingual use in enterprise deployments	Strong for long, polished operational writing	Cloud only; \$3 / \$15 per 1M	Cloud dependence; sparse explicit language documentation	42

Rank	Model	Why it ranks here	Language / corpus focus	Benchmarks / deployments	Access, licence, price	Key limitations	Sources
10	EuroLLM — European consortium, 2024	Included because many Antarctic programmes are European and sovereign/offline needs matter	Pan-European multilingual focus	Strong fit for multinational European stations or support orgs	Open weights	Less suitable for East Asian station languages than Qwen3 / Gemma / Command A	29

Antarctica verdict: the practical winners are **Command A**, **Qwen3**, and **Gemma 3** because they combine multilingual coverage with realistic self-hosting or local deployment. Pure cloud models remain useful, but their dependence on connectivity is a material weakness in Antarctic operations. 61

1 3 17 AfroBench: How Good are Large Language Models on ...

https://aclanthology.org/2025.findings-acl.976/?utm_source=chatgpt.com

2 6 7 Mansa

https://www.africanlanguageslab.com/mansa?utm_source=chatgpt.com

4 Jacaranda launches first-in-kind Swahili Large Language ...

https://jacarandahealth.org/jacaranda-launches-first-in-kind-swahili-large-language-model/?utm_source=chatgpt.com

5 40 How can I use the OpenAI API with text in different ...

https://help.openai.com/en/articles/6742369-how-can-i-use-the-openai-api-with-text-in-different-languages?utm_source=chatgpt.com

8 Paper page - InkubaLM: A small language model for low- ...

https://huggingface.co/papers/2408.17024?utm_source=chatgpt.com

9 Jacaranda/AfroLlama_V1

https://huggingface.co/Jacaranda/AfroLlama_V1?utm_source=chatgpt.com

10 Jacaranda/UlizaLlama3

https://huggingface.co/Jacaranda/UlizaLlama3?utm_source=chatgpt.com

11 Jacaranda/YorubaLlama

https://huggingface.co/Jacaranda/YorubaLlama?utm_source=chatgpt.com

12 Jacaranda/Xhosa_ZuluLlama3_v1

https://huggingface.co/Jacaranda/Xhosa_ZuluLlama3_v1?utm_source=chatgpt.com

13 inceptionai/jais-30b-v3

https://huggingface.co/inceptionai/jais-30b-v3?utm_source=chatgpt.com

14 silma-ai/silma-tts

https://huggingface.co/silma-ai/silma-tts?utm_source=chatgpt.com

15 Aya Expanse

https://docs.cohere.com/docs/aya-expanse?utm_source=chatgpt.com

- 16 43 55 59 **Google models | Gemini Enterprise Agent Platform**
https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/google-models?utm_source=chatgpt.com
- 18 **A Multi-task Indic Language Understanding Benchmark**
https://aclanthology.org/2025.naacl-long.507/?utm_source=chatgpt.com
- 19 27 44 **Qwen3: Think Deeper, Act Faster | Qwen**
<https://qwenlm.github.io/blog/qwen3/>
- 20 **SEA-LION: Southeast Asian Languages in One Network**
https://arxiv.org/abs/2504.05747?utm_source=chatgpt.com
- 21 **SILMA TTS: A Lightweight Open Bilingual Text to Speech ...**
https://huggingface.co/blog/silma-ai/opensource-arabic-english-text-to-speech-model?utm_source=chatgpt.com
- 22 **HyperCLOVA X Technical Report**
https://arxiv.org/html/2404.01954v1?utm_source=chatgpt.com
- 23 **cyberagent/calm3-22b-chat**
https://huggingface.co/cyberagent/calm3-22b-chat?utm_source=chatgpt.com
- 24 **krutrim-ai-labs/Krutrim-2-instruct**
https://huggingface.co/krutrim-ai-labs/Krutrim-2-instruct?utm_source=chatgpt.com
- 25 **sarvamai/OpenHathi-7B-Hi-v0.1-Base**
https://huggingface.co/sarvamai/OpenHathi-7B-Hi-v0.1-Base?utm_source=chatgpt.com
- 26 **MBZUAI/Llama-3-Nanda-10B-Chat**
https://huggingface.co/MBZUAI/Llama-3-Nanda-10B-Chat?utm_source=chatgpt.com
- 28 **Introducing Pharia-1-LLM: transparent and compliant**
https://aleph-alpha.com/en/blog/introducing-pharia-1-llm-transparent-and-compliant/?utm_source=chatgpt.com
- 29 39 **croissantllm/CroissantLLMBase**
https://huggingface.co/croissantllm/CroissantLLMBase?utm_source=chatgpt.com
- 30 **OpenLLM-France/Lucie-7B-Instruct-v1.1**
https://huggingface.co/OpenLLM-France/Lucie-7B-Instruct-v1.1?utm_source=chatgpt.com
- 31 **cortexso/sailor-2**
https://huggingface.co/cortexso/sailor-2?utm_source=chatgpt.com
- 32 **Overview: What is Pharia-1-LLM-7B?**
https://docs.aleph-alpha.com/phariaai-dev-guide/latest/pharia-llm/intro.html?utm_source=chatgpt.com
- 33 **LumiOpen/Poro-34B**
https://huggingface.co/LumiOpen/Poro-34B?utm_source=chatgpt.com
- 34 **INSAIT-Institute/BgGPT-Gemma-3-27B-IT**
https://huggingface.co/INSAIT-Institute/BgGPT-Gemma-3-27B-IT?utm_source=chatgpt.com
- 35 **OpenLLM-France/Claire-7B-FR-Instruct-0.1**
https://huggingface.co/OpenLLM-France/Claire-7B-FR-Instruct-0.1?utm_source=chatgpt.com
- 36 52 **google/gemma-3-12b-it**
https://huggingface.co/google/gemma-3-12b-it?utm_source=chatgpt.com
- 37 61 **Command A | Cohere**
<https://docs.cohere.com/docs/command-a>
- 38 **meta-llama/Llama-3.3-70B-Instruct**
https://huggingface.co/meta-llama/Llama-3.3-70B-Instruct?utm_source=chatgpt.com

41 **GPT-4.1 Model | OpenAI API**

https://developers.openai.com/api/docs/models/gpt-4.1?utm_source=chatgpt.com

42 **Introducing Claude 4**

https://www.anthropic.com/news/claude-4?utm_source=chatgpt.com

45 **Mistral Small 3.1**

https://mistral.ai/news/mistral-small-3-1/?utm_source=chatgpt.com

46 54 57 **Live API capabilities guide | Gemini API | Google AI for Developers**

<https://ai.google.dev/gemini-api/docs/live-api/capabilities>

47 **Can small language models revitalize Indigenous ...**

https://www.brookings.edu/articles/can-small-language-models-revitalize-indigenous-languages/?utm_source=chatgpt.com

48 49 53 **Latam-GPT: la primera IA regional abierta creada con ...**

https://cenia.cl/2026/02/10/latam-gpt-la-primera-ia-regional-abierta-creada-con-datos-latinoamericanos/?utm_source=chatgpt.com

50 **Sabiá-4**

https://www.maritaca.ai/en/blog/sabia-4/?utm_source=chatgpt.com

51 **Building Tucano 2: Open-Source Language Models That ...**

https://huggingface.co/blog/Polygl0t/tucano2?utm_source=chatgpt.com

56 **Llama 3.2: Revolutionizing edge AI and vision with open ...**

https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/?utm_source=chatgpt.com

58 **Antarctica**

https://education.nationalgeographic.org/resource/antarctica/?utm_source=chatgpt.com

60 **Introducing Llama 3.1: Our most capable models to date**

https://ai.meta.com/blog/meta-llama-3-1/?utm_source=chatgpt.com